The Binding Constraint in Layered Coordination Architectures

Sinéad O'Sullivan

Consider that network capability in coalition air operations has improved by an order of magnitude since Desert Storm. Similarly, sensor fusion is near-complete and data links are real-time. Meanwhile, the United States has spent hundreds of billions on JADC2, Link 16, and mesh networking.

Yet, friendly-fire rates have not declined and decision timelines have compressed beyond what the architecture can reliably process. ***Why?***

Because the problem is constrained by the weakness of the weakest layer, not augmented by the strength of the best performing one.

"Coordination architecture" is not a single capability but a *stack*: network transport at the bottom, then sensor fusion, identification, doctrinal harmonization, engagement authority, and coalition procedures at the top. System performance is therefore bounded by the weakest layer (the *binding constraint*) not the strongest.

A perfectly connected force that cannot agree on engagement authority in fifteen seconds has a coordination failure no network investment can fix.

To test this, I applied the Architecture Readiness Index, developed in my prior work on premature markets and coordination architectures, to the defense coordination stack. Each layer is scored using observable indicators derived from campaign data across six coalition air operations from 1991 to 2026. The top three layers are fully scored from published sources. The bottom three are partially scored, with methodology notes below.

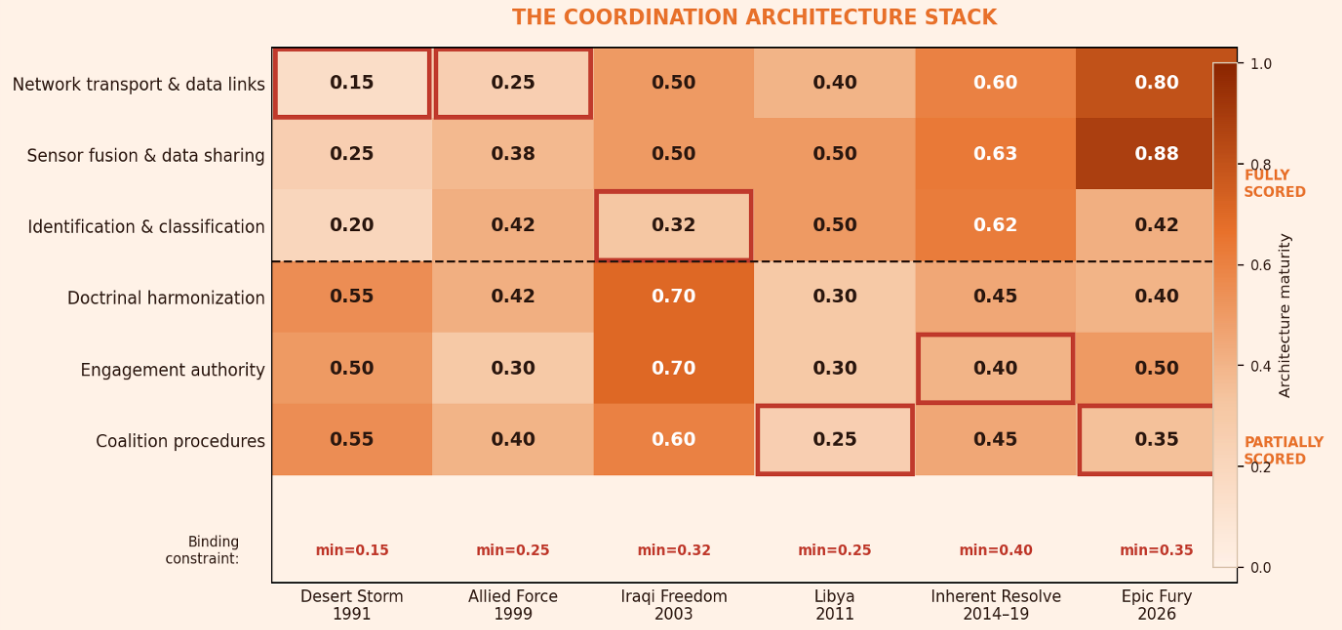

Figure 1. Architecture maturity by layer and campaign. Technical layers (above dashed line) show monotonic improvement. Institutional layers show no trend, and in some campaigns score lower than in 1991. Iraqi Freedom 2003 scores highest on institutional layers because it was a small, tight coalition with years of no-fly-zone experience. Red borders mark the binding constraint for each campaign.

Despite that network transport improved from 0.15 to 0.80 over thirty-five years, the institutional layers (doctrine, engagement authority, coalition procedures) are volatile and show no upward trend. Indeed, they are mission-specific, driven by coalition composition and political context rather than accumulated investment.

Thus, the binding constraint migrated upward as the technical layers improved.

Epic Fury (2026) is the critical case. It has the most advanced network architecture in the dataset (0.80) and the most capable sensor fusion (0.88). Yet its fratricide rate (2.0 per 10,000 sorties) is higher than every campaign except Desert Storm. The reason is visible in the heatmap: its weakest layer, coalition procedures, sits at 0.35, well below the viability threshold of 0.50, while decision time has compressed to fifteen seconds; twenty times more demanding than Desert Storm's 300-second window.

This generates a direct question: if the binding constraint is institutional, which interventions actually move it?

SENSITIVITY ANALYSIS

The sensitivity analysis tests four interventions against the Epic Fury baseline. The question is not which dial improves a layer score, as most of them do. The question is which dial *moves the binding constraint*, and therefore system output.

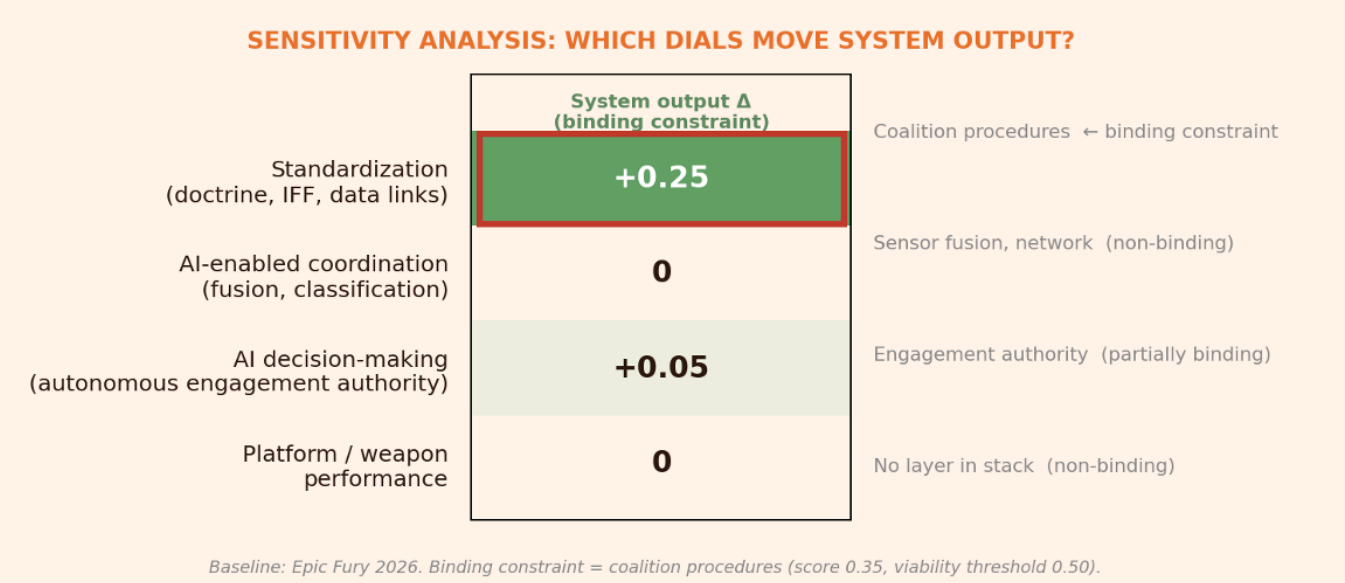

Figure 2. Sensitivity analysis: system output improvement (Δ binding constraint) from four intervention types, Epic Fury 2026 baseline. Viability threshold = 0.50. Only standardization addresses the binding constraint layer directly.

The answer is unambiguous:

- Standardization (of doctrine, identification protocols, and data link procedures) is the only intervention that directly addresses coalition procedures, the binding constraint layer. It produces a system output improvement of +0.25.
- AI-enabled coordination improves sensor fusion and network performance, both already non-binding, and produces zero system-level return.
- AI decision-making partially addresses engagement authority, a partially binding layer, for a modest +0.05.
- Platform and weapon performance improve no layer in the coordination stack and produce nothing.

This is not an argument against AI or platform investment, but an argument to highlight the need of *sequencing*.

Investment in non-binding layers before the binding constraint is resolved produces capability at the node with no improvement at the system level. The architecture is only as fast as its slowest institutional layer.

This finding generalizes beyond coalition air operations. Any layered coordination architecture exhibits the same structural property: technical layers improve through investment; institutional layers don't respond to the same logic.

As technical maturity accumulates, the binding constraint migrates into the institutional stack, where it is invisible to standard investment metrics and immune to standard investment tools.

The implication is direct: before allocating resources to any layer, identify which layer is binding.

SCORING METHODOLOGY

Each layer is scored on a 0–1 scale using the indicators below. Scores for the top three layers are derived directly from published data. Scores for the bottom three are constructed from partial data with explicit assumptions noted. All scores are author's calculations from the listed sources.

Layer	Indicator	Source & method
Network transport	% coalition platforms on shared real-time data link (JTIDS/Link 16/mesh)	CRS campaign reports; DISA interoperability assessments. Fully scored.
Sensor fusion	Cross-domain sensor types in common operating picture, normalized to max (8)	CENTCOM AARs; ISR integration assessments. Fully scored.
Identification	IFF mode generation × (1 – normalized fratricide rate per 10K sorties)	CSIS; UK MoD Board of Inquiry; CENTCOM. Fully scored.
Doctrinal harmonization	% coalition partners under unified doctrine, weighted by sortie share	NATO JAPCC; CRS coalition analyses; UK MoD. Partially scored.
Engagement authority	Kill-chain delegation level × (1 – proportion of targets requiring national approval)	CENTCOM AARs; NATO assessments. Partially scored.
Coalition procedures	Pre-campaign exercise duration + interoperability certification level	RAND; NATO exercise records. Partially scored.

KEY DATA INPUTS

Sortie counts: Desert Storm ~116,000 (CRS); Allied Force ~38,000 (CRS); Iraqi Freedom ~41,000 (CENTCOM); Libya ~26,000 (NATO); Inherent Resolve ~110,000 (CENTCOM cumulative); Epic Fury ~15,000 (estimated from UK MoD and CENTCOM reporting). Friendly-fire incidents from CENTCOM AARs and UK MoD Board of Inquiry reports.

Viability threshold: Set at 0.50, based on the ARI criterion that system-level coordination requires all layers above a minimum functional level. No campaign in the dataset achieves a minimum layer score above 0.50. The threshold is consistent with the formal viability condition $\Delta(t) \leq \theta$ from the companion papers.

This article applies the Architecture Readiness Index from two working papers by Sinéad O'Sullivan: Institutions as Coordination Architectures and Market Formation as a Systems Engineering Problem. Copies available on request: s@sinead.co.